

PRODUCT PERSONAS 1



Demographics

- Name: Alex Chen
- Age: 32
- Location: Seattle, USA
- Job: Senior Backend / Platform Engineer
- Salary: \$120,000–\$150,000/year
- Experience: 7+ years in backend and distributed systems
- Team: Platform engineering team supporting internal services

Profile

- Experienced backend engineer managing infrastructure for multiple services.
- Works on integrating AI capabilities into production systems.
- Prefers reliable, scalable, and modular architectures.
- Focuses on maintaining stable APIs for internal development teams.
- Seeks centralized solutions for managing multiple LLM providers.

Interests

- Scalable backend system design
- Distributed systems and fault tolerance
- Observability tools (logs, metrics, tracing)
- Cloud and container technologies
- Performance and cost optimization

Frustrations

- Different LLM providers have inconsistent APIs.
- Service downtime affects dependent applications.
- Difficult to track token usage and costs.
- Lack of centralized monitoring for AI requests.
- Switching providers requires code changes in multiple services.

Goals

- **Enable vendor-agnostic LLM integration.**
- **Maintain high reliability and uptime.**
- **Provide a stable API for internal teams.**
- **Improve monitoring and observability.**
- **Track AI usage and cost efficiently.**

How Resilient LLM Gateway Helps

- Automatic Routing & Failover – switches providers if one fails.
- Unified API – same interface for all LLM providers.
- Structured Logging – better debugging and monitoring.
- Usage Tracking – monitor token usage and cost.
- Provider Abstraction – change providers without modifying apps.

